

A golden Centaur archer is depicted in a dynamic pose, holding a bow and arrow, set against a dark purple space background filled with stars and glowing energy fields.

# CHIEON

*By @rodsoto @josephzadeh*

*\$whoami...*

*Rod Soto*

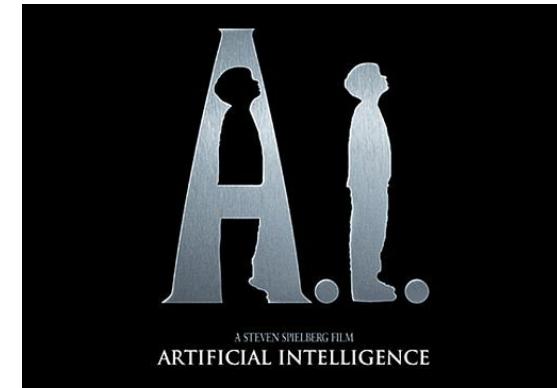
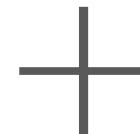
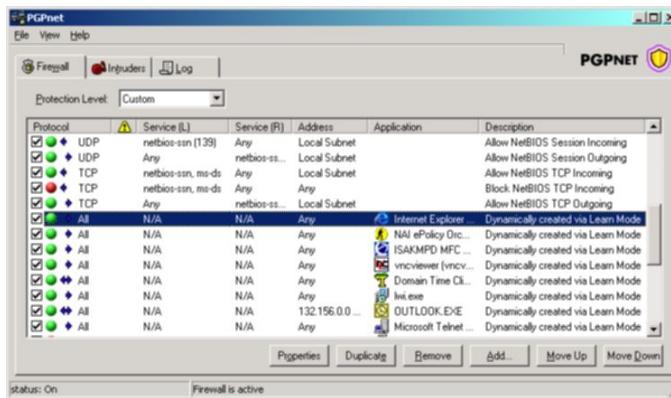
Director of Security Research at JASK.AI. Co-founder of Hack The Valley %27. Creator/Developer of Kommand & KonTroll CTF. Co-founder of Hackmiami. Previously at Prolexic, Caspida, AKAMAI, Splunk.

*Joseph Zadeh*

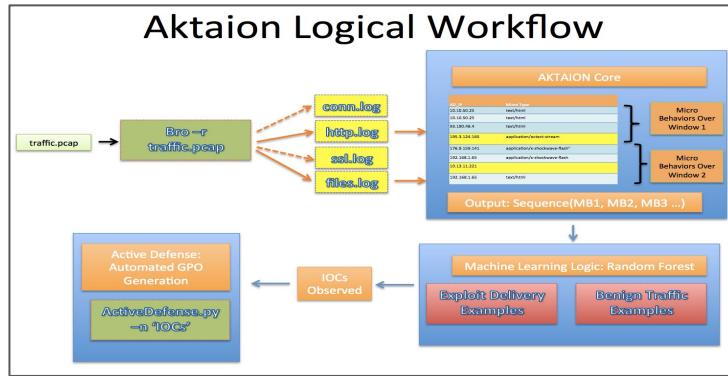
Director of Data Science at JASK.AI. Co-founder of Hack The Valley %27. Previously at Pivotal, Kaiser Permanente, Caspida, Splunk.

# *INTRODUCTION*

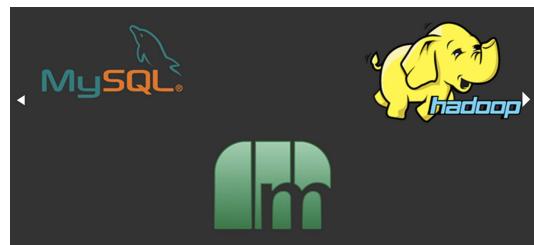
# ML + FW = Next Generation of Home Defense



# Motivation: Security Innovations are hard to scale down



**OpenSOC**  
Big Data Security  
Analytics Framework



# Challenges in home networking defense

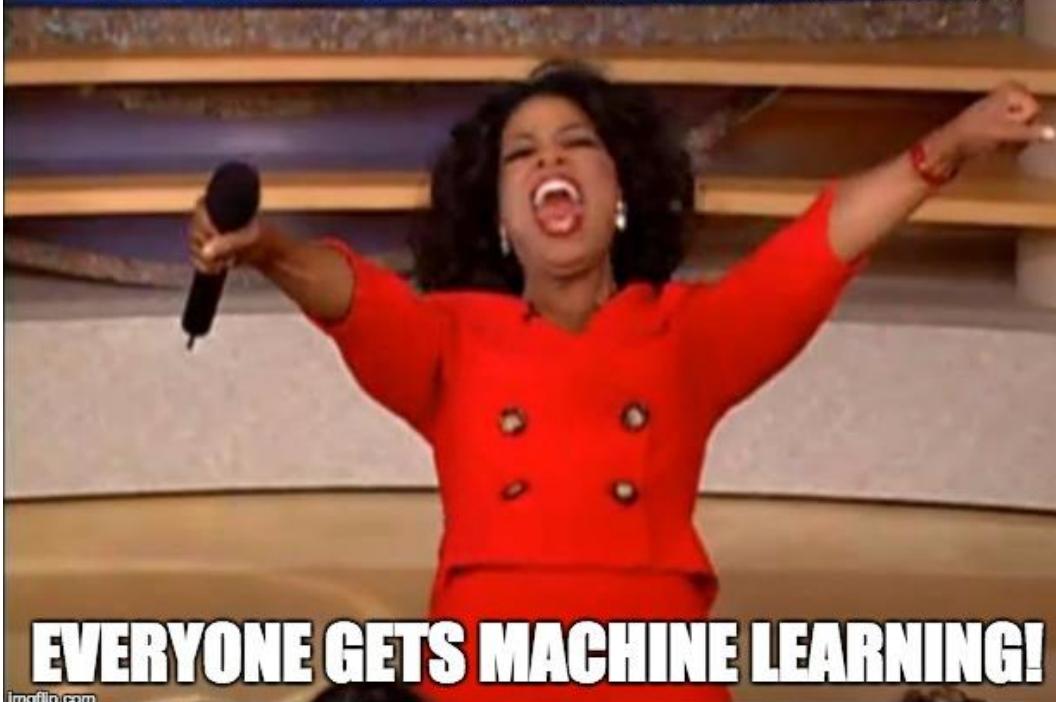
- Internet of Things (toasters, refrigerators, thermostats, cameras, wearables, AI, routers, printers, etc)
- Unpatched vulnerable devices become Zombies of very large botnets feeding the crime ecosystem (CPEs are prime target)
- Common users have no idea or knowledge of what to do with all these “home” infrastructure becoming prey of snake oil pharisees
- Huge privacy risks, (I.E Unknown external IP accessing webcam or child monitor, somebody unlocking doors remotely, large data transfer from Home NAS to unknown IP address)

# Challenges in home network defense II

- Wifi poachers, p2p piracy, someone using your home network as “VPN”
- Targeted attacks against your organization via your home network
- Targeted attacks against personal property (identity, finance, insurance)
- Port 25/110/143/993/465 running?. Is your home network being used for SPAM
- 9001/9030/9050/9150 ? Why is somebody using TOR at home?

# BRINGING MACHINE LEARNING TO HOME DEFENSE

**YOU GET MACHINE LEARNING!**



**EVERYONE GETS MACHINE LEARNING!**

# No longer a silicon valley, tech class thing

- Our effort is to bring a basic analytics tool to the masses. We believe this technology can help everyone, not just large enterprises.
- We tried to make it simple and easy to use.
- Machine learning can help everyone at home and in the SMB discover certain classes of threat actors
- Seeing is believing, these basic statistics will provide an overall yet sufficient picture to understand what is going on at your home network

Let's be practical about what software automation can uncover. ML or any single point solution is not good for defending against advanced actors.

# Advantages of ML

- Using ML allows us to put together simple analytics, on the main elements of home networking: First class citizens in our data model (Assets, Services, Protocols, Users, Traffic).
- ML allows us to go beyond of static signature based technologies. This aims to augment what we are good at so far and apply evidence fusion to point solution output like Bro/Snort alerts or iptable type log changes (osquery).
- IDS best practices + analytic algorithms help creates a scenario where detection of threats based on dynamic and multi contextual indicators is possible.

# CHIRON - ML network home security use cases

- Malicious file downloads
- Fingerprints users, services, and protocols
- High risk users, assets, domains
- Data usage
- North - South traffic
- Relationships between assets
- Social Media usage
- Odd application/traffic/services

# Bringing Machine Learning to home defense

- You don't need a hadoop cluster at home in order to use machine learning
- There are several open source tools that we can use to have a decent proof of concept of network analytics plus machine learning
- No “SNAKE OIL” OPEN SOURCE [github.com/jzadeh/chiron](https://github.com/jzadeh/chiron)
- No need for a network tap, we are using passive detection tool pof in combination with NMAP. You can use a network tap and pipe the output if you want.
- It can be run from security onion
- FIRE AND FORGET NO NEED TO CODE

# A case to bring ML to home defense

- Home networks are your castle
- Pivot of targeting/phishing employees at their homes (remote workers)
- IoT -- big leakers of private info, unpatch, vulnerable, significant risk
- Do you live in highly dense building? Anybody poaching your Internet service?
- Where all those devices connect to?
- Where all my users connecting to?
- Is there any suspicious NORTH-SOUTH traffic?
- Dynamic asset discovery

# A case to bring ML to home defense

Three basic tasks:

- Perform basic discovery and analytics of home network assets (IoTs, Workstations, Laptops, Servers, routers)
- Fingerprints users, services, and protocols
- Applies analytics to users and devices (Average session length, Traffic, Visited sites)

# The Tool

IDS Source Code: <http://www.github.com/jzadeh/chiron-elk>

VM ISO: [https://drive.google.com/uc?id=oB8Goyrax\\_wtaNzAxM2ptdUZ6anc&export=download](https://drive.google.com/uc?id=oB8Goyrax_wtaNzAxM2ptdUZ6anc&export=download)

## Requirements

- Can be deployed on a VM - 60GB HD 1GB RAM
- Chiron is good for a home network 5 - 10 users
- 30 day log retention standard
- ELK for visualization and simple interactive analysis

# Core Principles

- **Lambda defense:** Use of big data tech for distributed processing and analytics
- **Micro behavior:** Basic units within a exploit chain or threat
- **Deployment Model:** Target a specific home use demographic for ease of use detection of change point behaviors linked to threats.

# The Lambda defense at home...

- Large organizations are currently attempting to scale intrusion detection by using big data distributed processing technologies (mainly Hadoop).
- CHIRON does not depend on a Hadoop infrastructure even though it does use some libraries part of the ecosystem. CHIRON is not meant to scale to 1000's of users however it is based on the same analytic principles and some of its ML frameworks
- By identifying, analyzing and providing key analytics for home network security, CHIRON gives home defense insights that is simply not available at this point without technical expertise and hardware.

# Micro behaviors what is that?

- Basic units present within a context of a threat or exploit kill chain. In order to define Micro behaviors we need to label data and give it a value within a context of threat or exploit kill chain.
- List of Micro behaviors present in a ransomware attack:
  - Payload delivery.
  - Focused on traffic to malicious sites and the related indicators when malicious code is served.
  - Including things such as URI entropy, redirects, domain generated by algorithms (DGAs), types and sequences of MIME content presented to victim during payload delivery.

# Micro behaviors what is that? II

- Call back (Phone home) patterns, including user agent, URI strings, HTTP “GET” or “POST” requests, DNS queries, URI strings, frequency of call-backs, periodicity of connections.
- Covert Channel indicators, such as non HTTP traffic (HTTPS), and non DNS traffic present during such communications.
- Entropy in URI/URL (number of dots, special characters, URL crazy) - Whois (registrar with bad rep / geolocation / bulletproof hosting

# What are Micro behaviors? III

- URL Shorteners (Tiny URL, bit.ly, etc)
- Spelling errors, passwords fields
- Presence of iframes, javascript, pagerank?, source of images - Presence of .exe .pdf .bat .ps1 .ps .bin .bat .jar .bin .zip

By analyzing these Micro behaviors within a context of exploitation and entity behavior we can develop methods for detection. Then train the models, then test detection.

# Microbehavior : A definition to express behaviors associated to TTP's in Code

```
package com.aktaion.ml.behaviors

import scala.util.matching.Regex

class ExploitationUriBehaviors extends MicroBehaviorSet {
    val uriMaxPathDepth = MicroBehaviorData("MaxPathDepth", "Maximum path length of all URI's observed")
    val uriMinPathDepth = MicroBehaviorData("MinPathDepth", "Minimum path length of all URI's observed")
    val uriMaxLength = MicroBehaviorData("MaxUriLength", "Maximum length of all URI's observed")
    val uriMinLength = MicroBehaviorData("MinUriLength", "Minimum length of all URI's observed")
    val uriDistinct = MicroBehaviorData("UniqueNumberOfUri", "Unique count of URI's in window")
    val uriMaxEntropy = MicroBehaviorData("UriMaxEntropy", "Maximum entropy of all URI's observed")
    val uriMinEntropy = MicroBehaviorData("UriMinEntropy", "Minimum entropy of all URI's observed")
    val base64Match = MicroBehaviorData("UriBase64", "URI's observed contain large number of base 64 strings")
    val percentEncodingMatch = MicroBehaviorData("UriBase64", "URI's observed contain large number of percent encoded strings")

    val behaviorVector = List(...)

    val encodingBase64 = new Regex( """^(?:[A-Za-z0-9+/]{4})*(?:[A-Za-z0-9+/]{2}==|[A-Za-z0-9+/]{3}=)?$""")
    val encodingBase64modified = new Regex( """^([A-Za-z0-9+/]{4})*([A-Za-z0-9+/]{4}|[A-Za-z0-9+/]{3}=|[A-Za-z0-9+/]{2}==)$""")
    val encodingPercentSimple = new Regex( """%[A-Fa-f0-9]{2}""")}

    class ExploitationTimingBehaviors extends MicroBehaviorSet {
        val maxTimeIntervalA = MicroBehaviorData("MaxTimeIntervalA", "Difference in timestamp between event 1 and event 2")
        val maxTimeIntervalB = MicroBehaviorData("MaxTimeIntervalB", "Difference in timestamp between event 2 and event 3")
        val maxTimeIntervalC = MicroBehaviorData("MaxTimeIntervalC", "Difference in timestamp between event 3 and event 4")
        val maxTimeIntervalD = MicroBehaviorData("MaxTimeIntervalD", "Difference in timestamp between event 4 and event 5")

        val minTimeIntervalA = MicroBehaviorData("MinTimeIntervalA", "Difference in timestamp between event 1 and event 2")
        val minTimeIntervalB = MicroBehaviorData("MinTimeIntervalB", "Difference in timestamp between event 2 and event 3")
        val minTimeIntervalC = MicroBehaviorData("MinTimeIntervalC", "Difference in timestamp between event 3 and event 4")
        val minTimeIntervalD = MicroBehaviorData("MinTimeIntervalD", "Difference in timestamp between event 4 and event 5")

        val intervalLength = MicroBehaviorData("TimeLengthInWindow", "Difference last and first time in the window")

        val ratioOfDeltasA = MicroBehaviorData("DeltaRatioA", "Count of time delta < 1 second over window size")
        val ratioOfDeltasB = MicroBehaviorData("DeltaRatioB", "Count of time delta < 5 second over window size")
        val ratioOfDeltasC = MicroBehaviorData("DeltaRatioC", "Count of time delta < 10 second over window size")
        val ratioOfDeltasD = MicroBehaviorData("DeltaRatioD", "Count of time delta < 20 second over window size")
        val ratioOfDeltasE = MicroBehaviorData("DeltaRatioE", "Count of time delta >= 100 second over window size")

        val behaviorVector = List(...)}
    }
```

**CHIRON ARCHITECTURE**

# Scaling Op-Sec to the Home Environment

- There are tools for defense but there is a learning curve to deployment models see for instance “DEFCON 17: Asymmetric Defense: How to Fight Off the NSA Red Team with Five People or Less”  
[https://www.defcon.org/images/defcon-17/dc-17-presentations/defcon-17-eistratios\\_gavas-nsa\\_red\\_team.pdf](https://www.defcon.org/images/defcon-17/dc-17-presentations/defcon-17-eistratios_gavas-nsa_red_team.pdf)

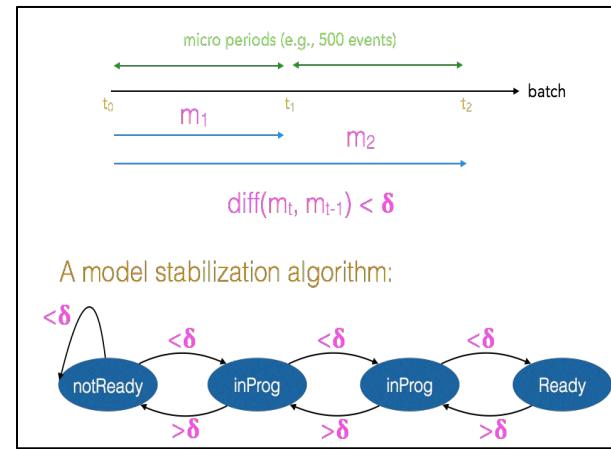
Some of the tools mentioned + bootstrapping FreeBSD

Monowall, IPCop, Untangle, pfSense

Goal is to make these type of techniques easier to adopt for non-power users

# How Does CHIRON Learn?

1. Key insight is to take action on new behaviors observed on a small network segment/s
2. New Behavior means a change point in underlying data:
  - a. New service interaction ( IP1  $\Rightarrow$  IP2 via port /app never observed before)
  - b. New DMZ interaction
  - c. Exploit chain observed with new MIME Type
3. Exploit detection learning (standalone model based off Aktaion)



# Adversary Modeling For The Home

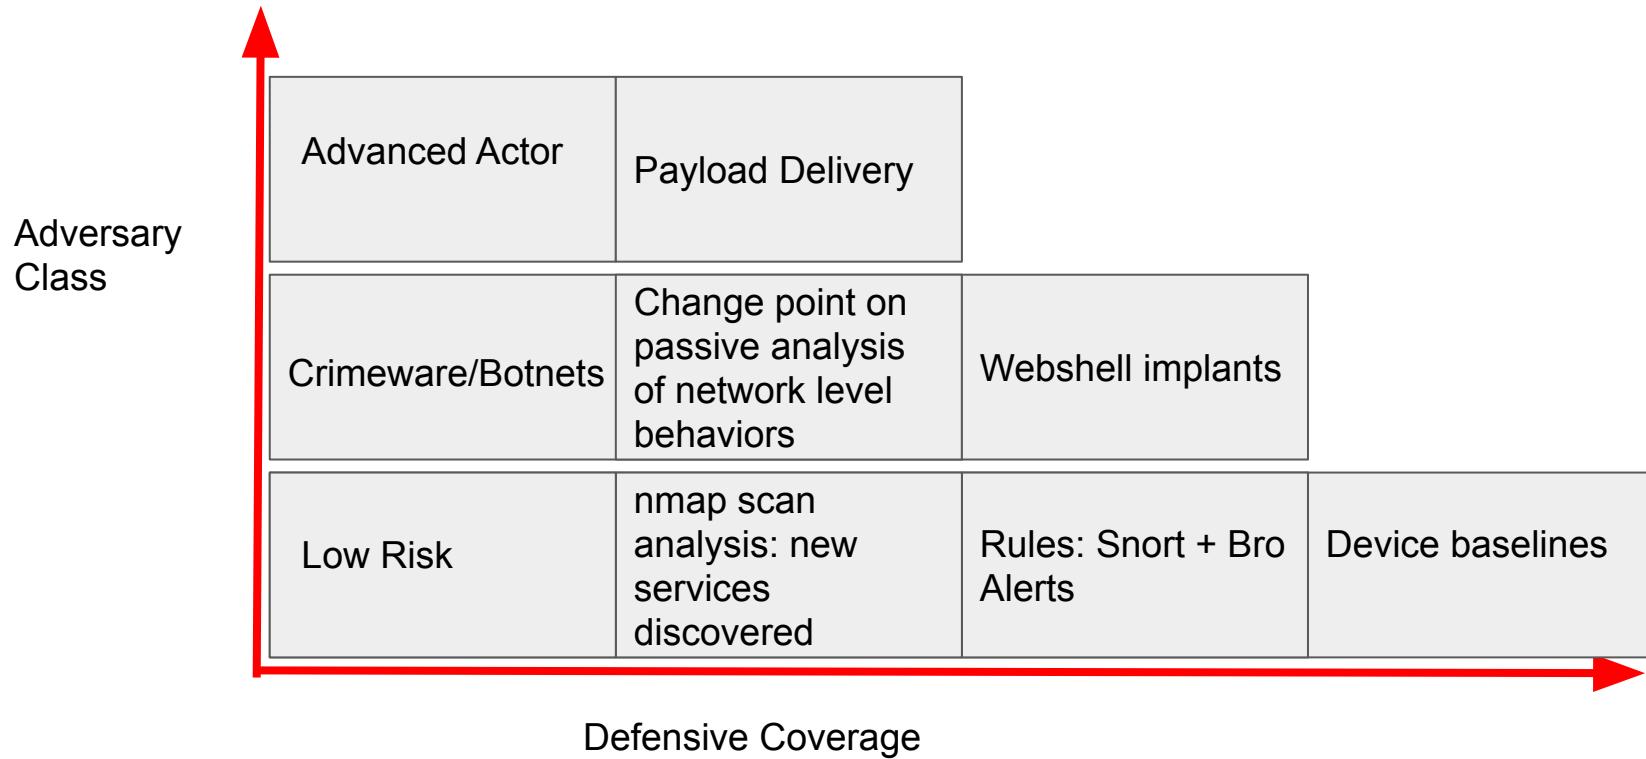
Key classes to model:

1. Crimeware / Ransomware
2. Opportunistic Neighbors/Local attacks
3. IoT/BYOD exploitation (android attacks)
4. Merai/Botnets

Key Use Cases to Targeted in V1:

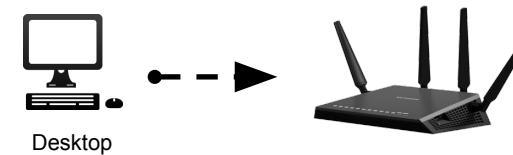
1. Delivery of payloads (weaponized code/RATs and implants on hosts)
2. IoT / Perimeter Device attacks

# Adversary Modeling For The Home/SMB

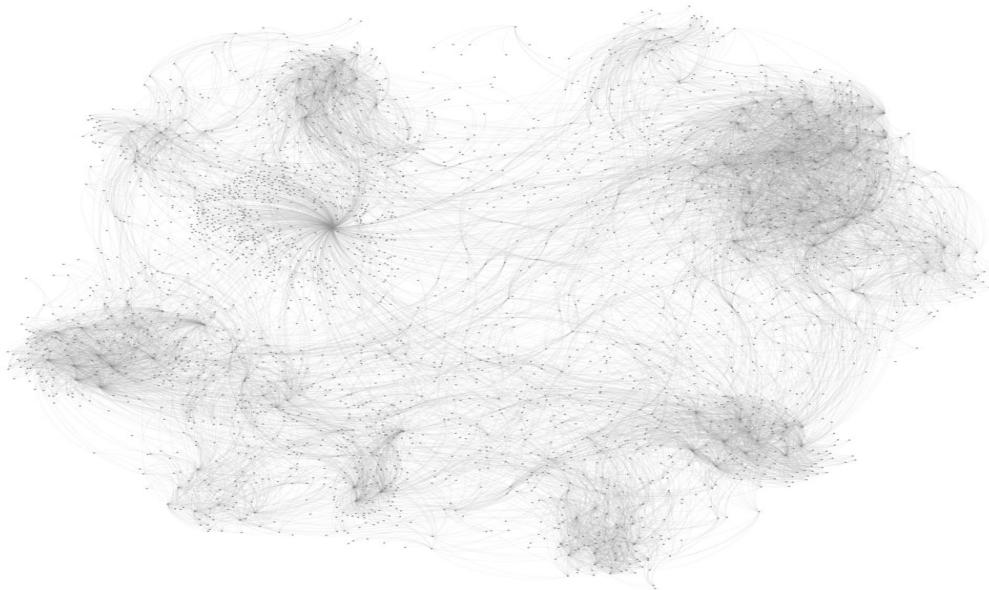


# Graph Behaviors For Home Defense

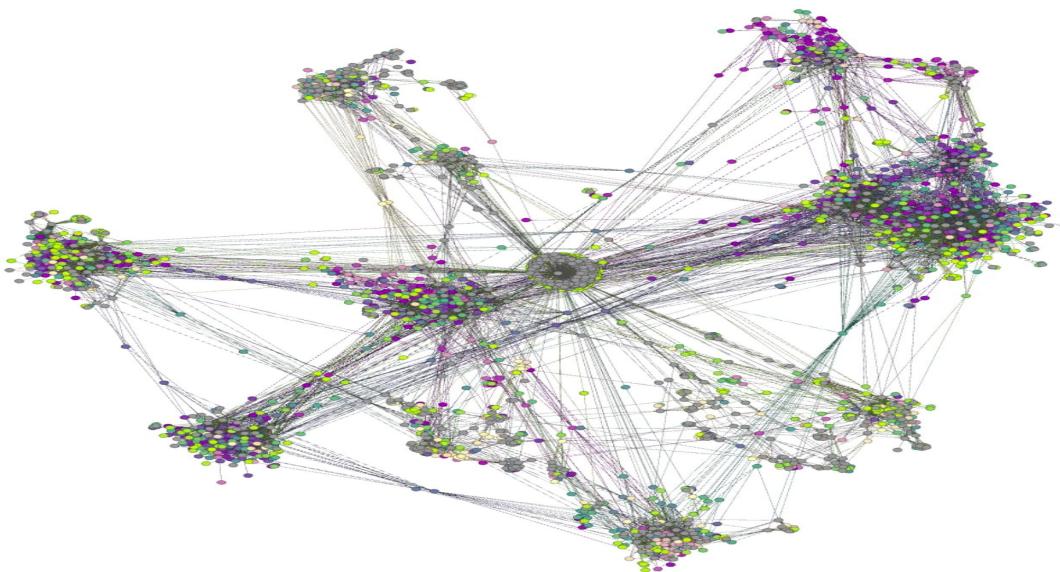
Key intuition: Anomalies on rare services and rare behaviors not in the 30 day log history



# Seeing the Analytic In Action

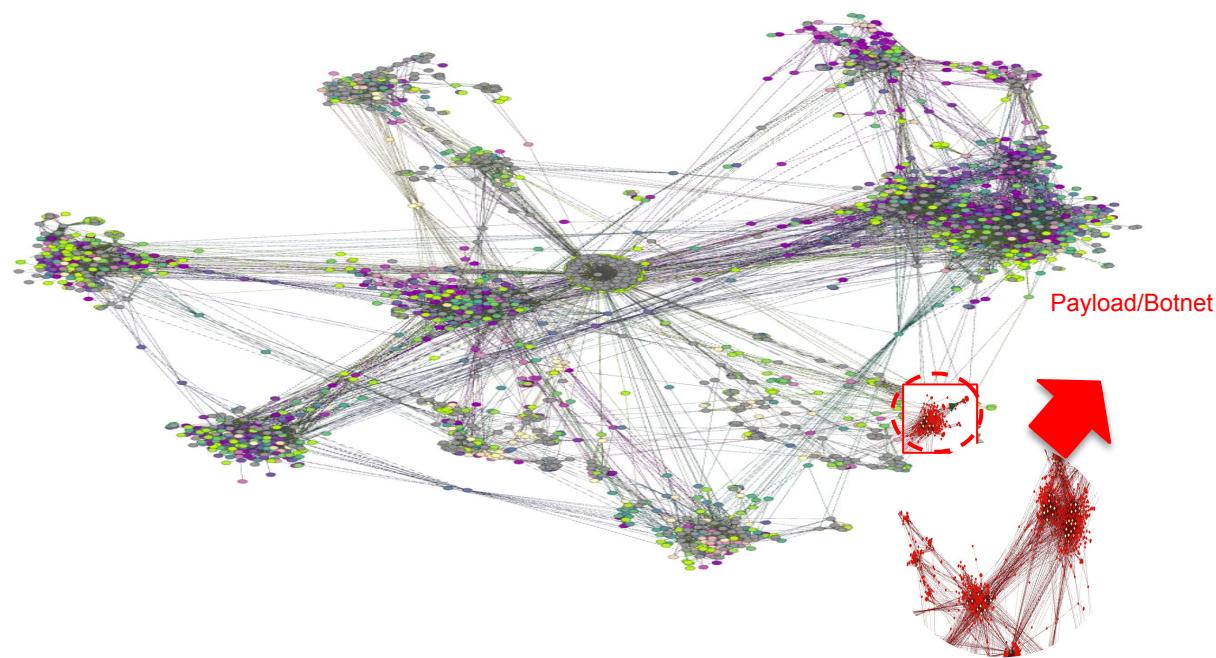


# Seeing the Analytic In Action

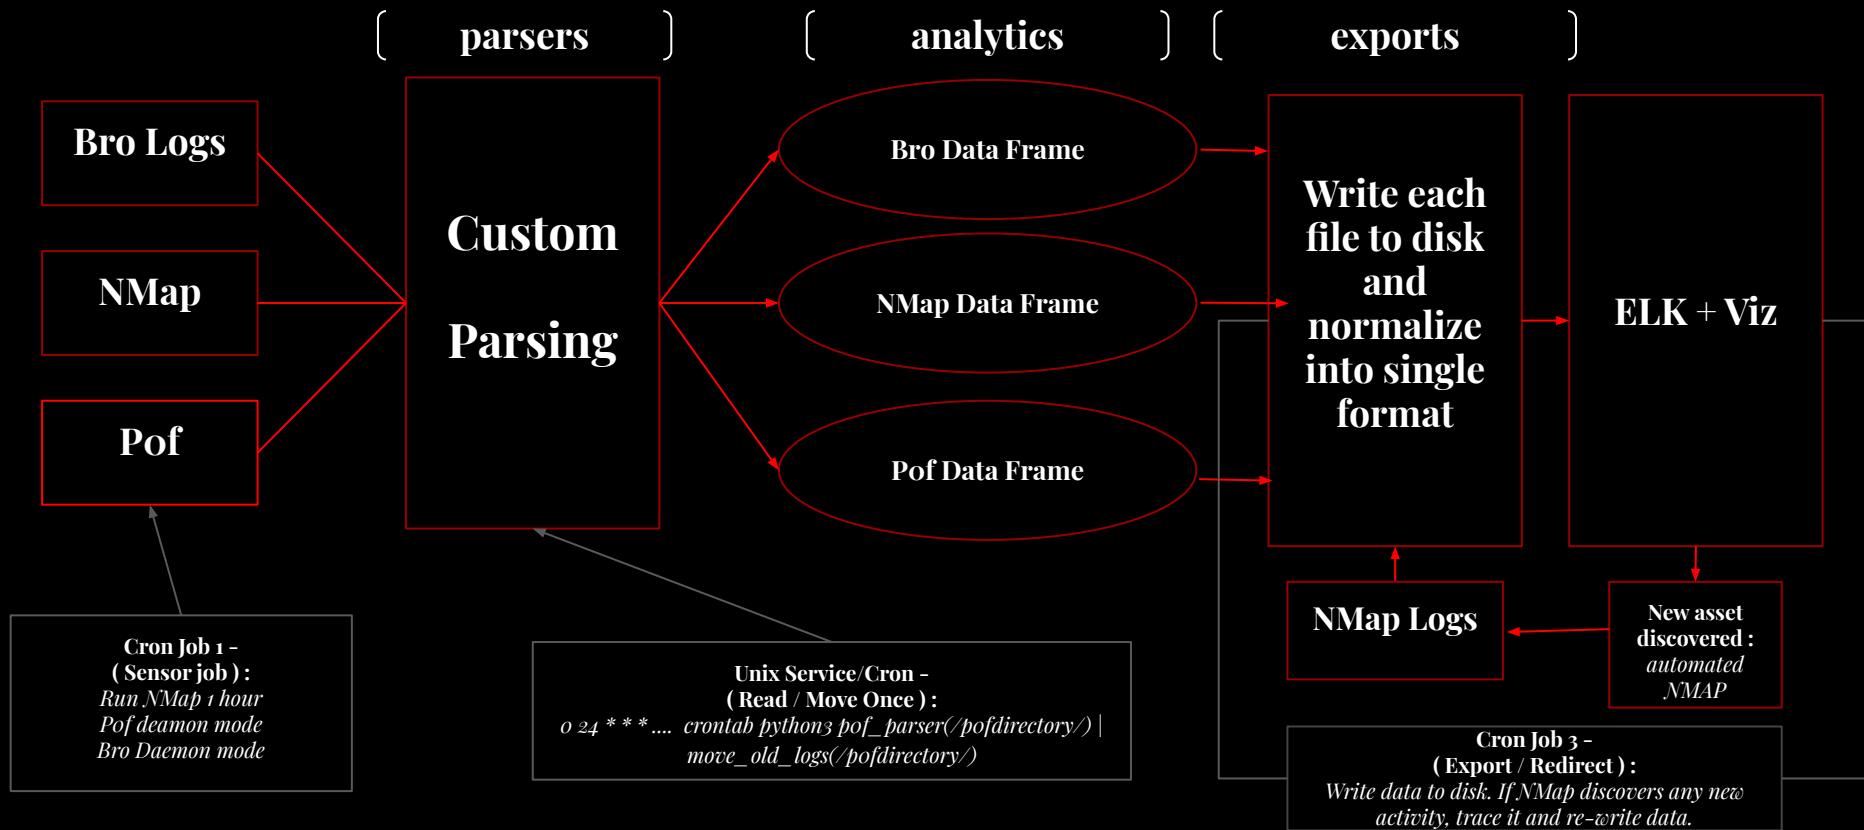


# Seeing the Analytic In Action

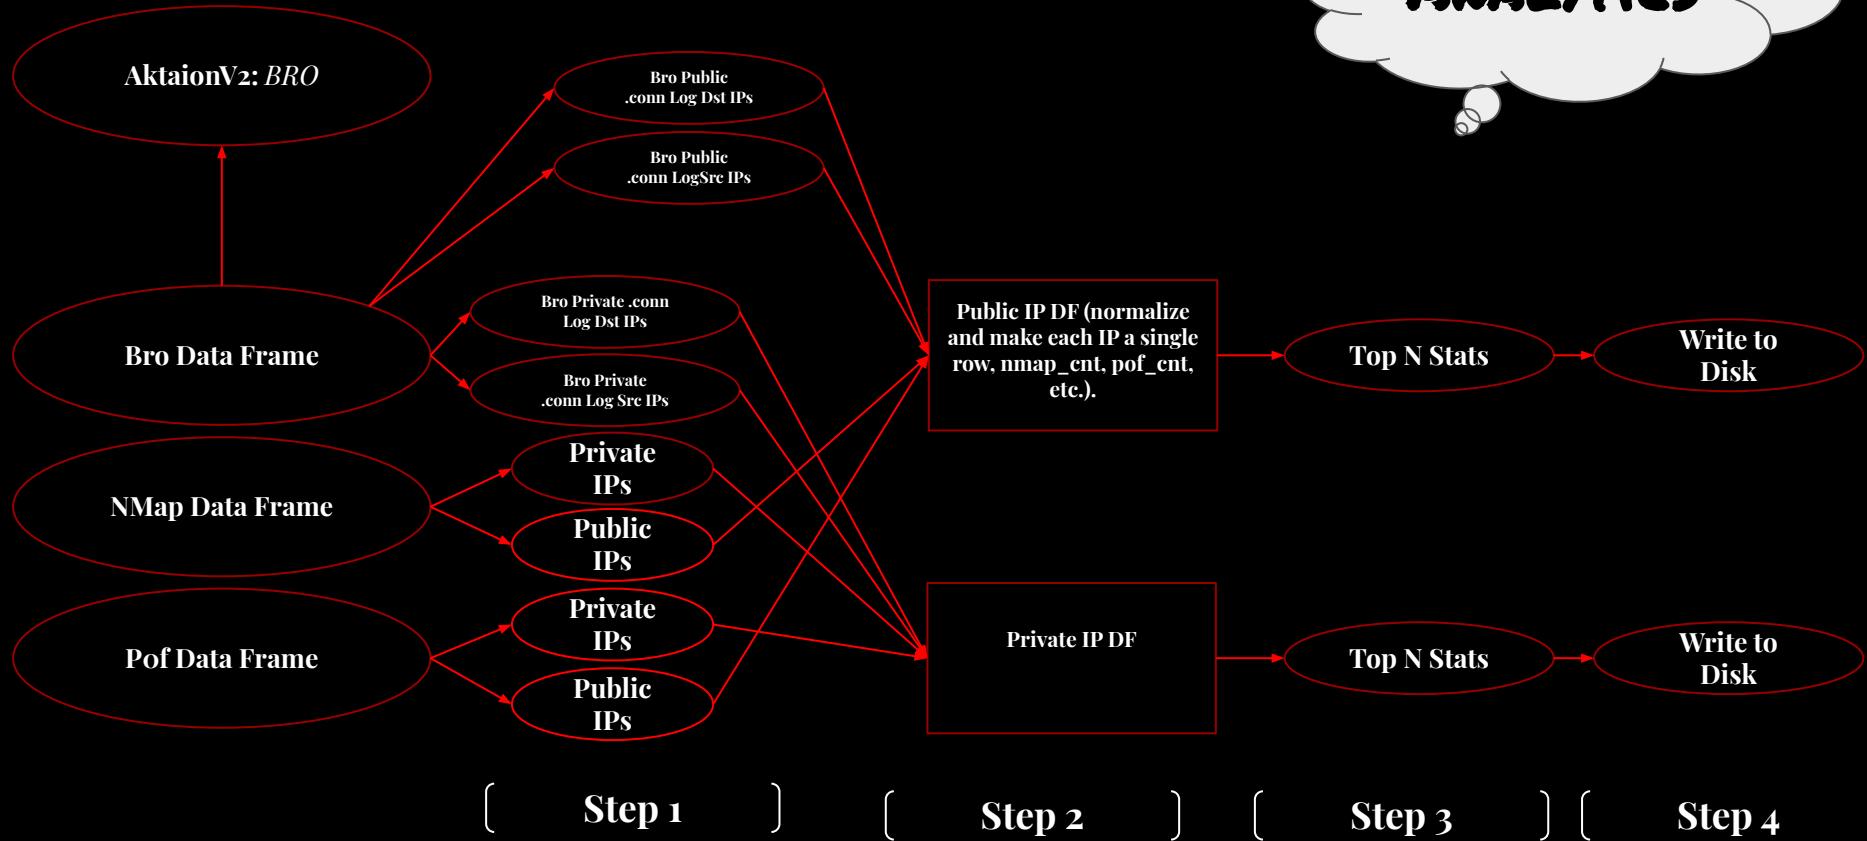
Changepoint detection at the service/application layer



# ML Exploit Detection: Aktaion V2



# ANALYTICS



# The Tool

IDS Source Code: <http://www.github.com/jzadeh/chiron-elk>

VM ISO: [https://drive.google.com/uc?id=0B8G0yrax\\_wtaNzAxM2ptdUZ6anc&export=download](https://drive.google.com/uc?id=0B8G0yrax_wtaNzAxM2ptdUZ6anc&export=download)

## Requirements

- Can be deployed on a VM - 60GB HD 1GB RAM
- Chiron is good for a home network 5 - 10 users
- 30 day log retention standard

# V2 of Architecture

Leverage OS query to do active defense interaction for targeting one specific use case: Given a network flow on the home Router what process on what host generated that flow.

Snort integration

Shadow IDS (active response disabled)

- Receive but no Transmit for a hardwired network tap

Community Involvement!

**EXPLOIT DETECTION ML:  
AKTAION**

# Phishing Attacks Had A Big Year In 2016, But Exploit Kits Will Make A Comeback In 2017



Spandas Lui

Feb 7, 2017, 12:30pm · Filed to:

Australian Stories ▾

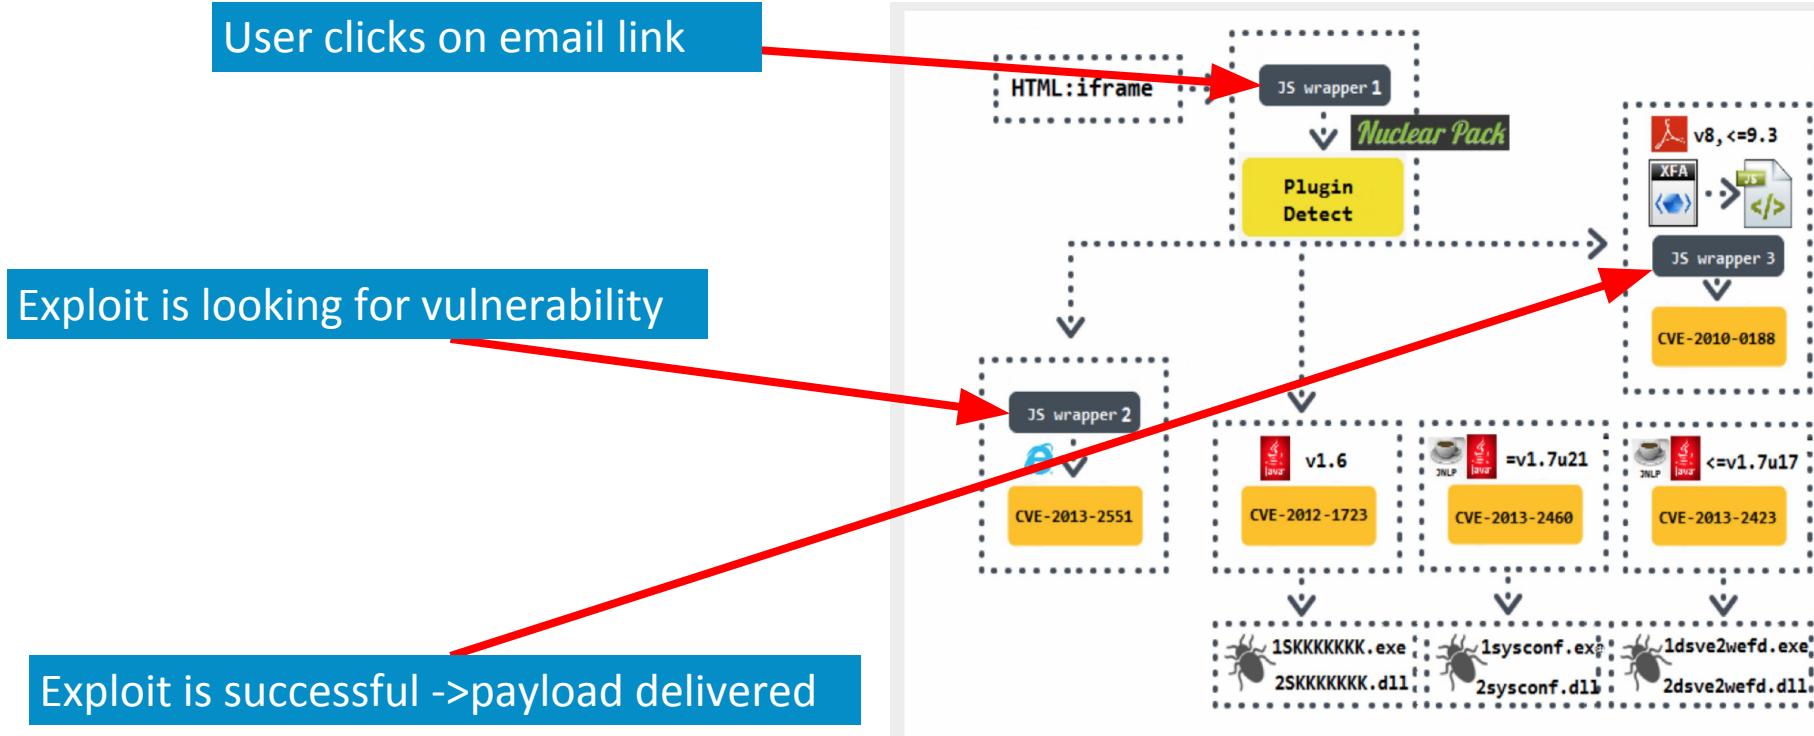
Share [f](#) [t](#) [in](#) [JU](#) [G](#)



Image: iStock

Last year saw a steep increase in malware distributed through email attachments in phishing attacks but 2017 will be the year of malware exploit kits, according to a report by security vendor Malwarebytes. We take a closer look at malware trends from 2016 and

# Successful Exploit Chain Post Phish (Angler EK)

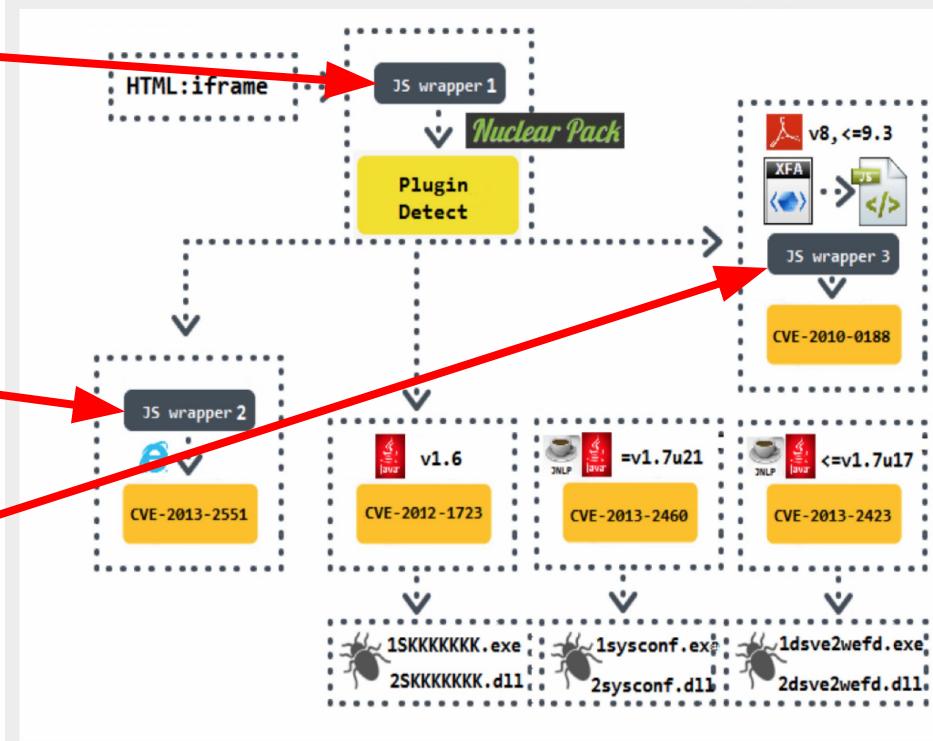


# This is what it looks like in web proxy logs

1. **Initial Redirect From Poisoned Domain:** [29/Apr/2015:16:52:23 -0700] "Nico Rosberg" 192.168.122.177 69.162.78.253 1500 200 TCP\_HIT "GET <http://forbes.com/gels-contrariness-domain-punchable/1.html/548828415920276748> HTTP/1.1" "Internet Services" "low risk" "text/html" 604 142 "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0)" "http://forbes.com/gels-contrariness-domain-punchable/1.html" "-" "0" "" "-"

2. **Flash Exploit:** [29/Apr/2015:16:52:26 -0700] "Nico Rosberg" 192.168.122.177 69.162.78.253 1500 200 TCP\_HIT "GET [http://portcallisesposturen.europartsplus.org/IMV0BZKPLqAJYIDe02t5hMMNyZBLN\\_q4kafJkVNqJVTnTmd](http://portcallisesposturen.europartsplus.org/IMV0BZKPLqAJYIDe02t5hMMNyZBLN_q4kafJkVNqJVTnTmd) HTTP/1.1" "Internet Services" "low risk" "application/x-shockwave-flash" 518 821 "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0)" "http://forbes.com/gels-contrariness-domain-punchable/1.html/548828415920276748" "-" "0" "" "-"

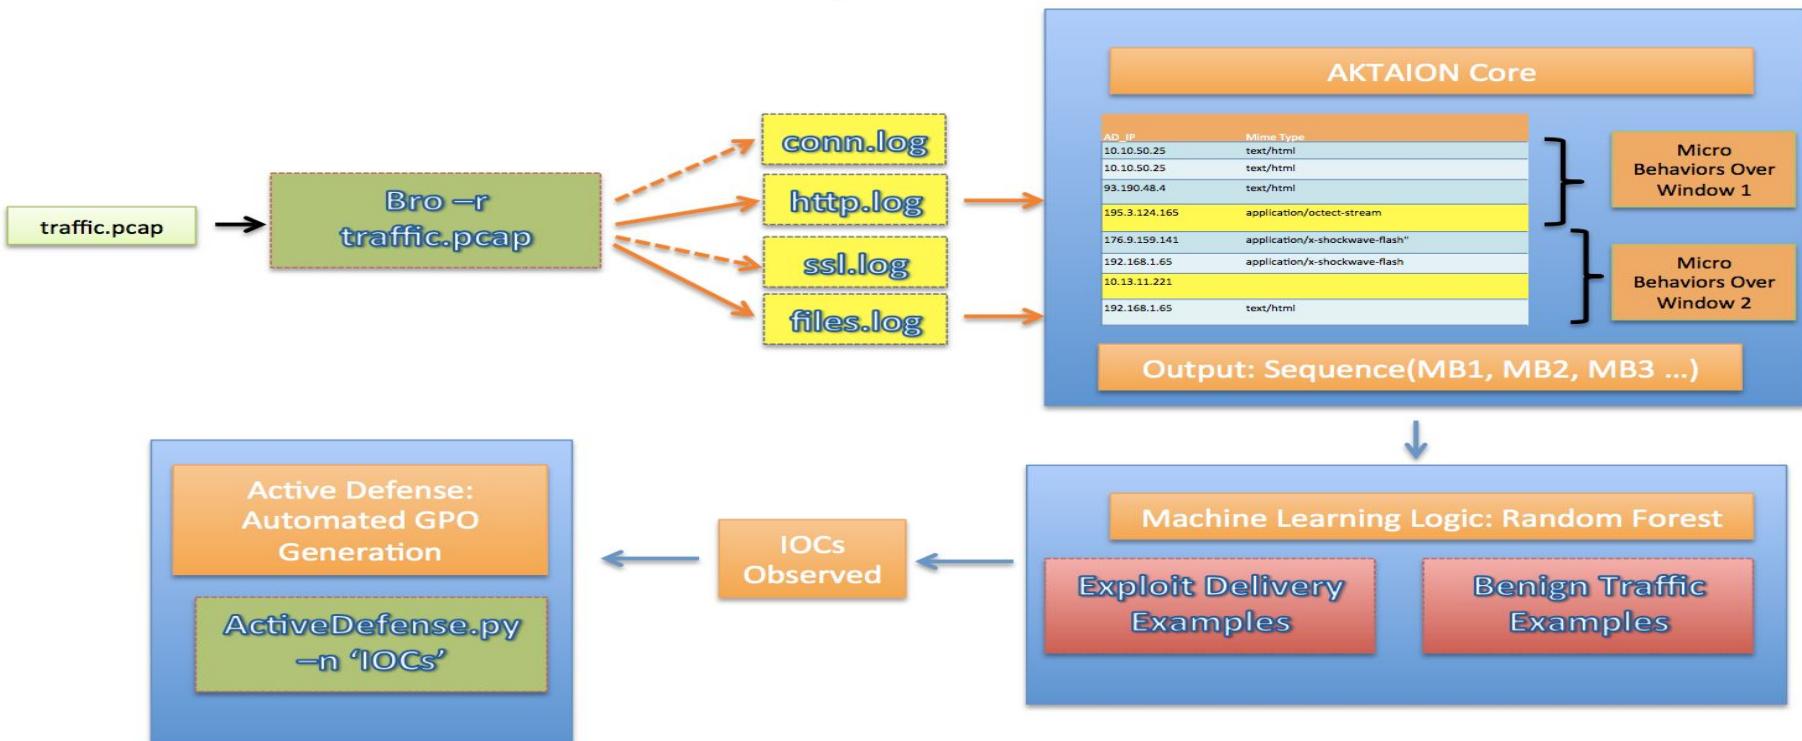
3. **Payload:** [29/Apr/2015:16:52:27 -0700] "Nico Rosberg" 192.168.122.177 69.162.78.253 1500 200 TCP\_HIT "GET [http://portcallisesposturen.europartsplus.org/UX7n1YkbNn8FUV6QVtEZLj\\_p\\_gLvRKIWEWmz3r7Ug8suRiY\\_](http://portcallisesposturen.europartsplus.org/UX7n1YkbNn8FUV6QVtEZLj_p_gLvRKIWEWmz3r7Ug8suRiY_) HTTP/1.1" "Internet Services" "low risk" "application/octet-stream" 136 915 "" "" "-" "0" "" "-"



# Workflow

1. Take PCAPs of known (labeled) exploits and known (labeled) benign behavior and convert them to bro format
2. Convert each Bro log to a sequence of micro behaviors (machine learning input)
3. Compare the sequence of micro behaviors to a set of known benign/malicious samples using a Random Forest Classifier (  
<http://weka.sourceforge.net/doc.dev/weka/classifiers/trees/RandomForest.html>)
4. Derive a list of indicators from any log predicted as malicious
5. Pass the list of IOCs (JSON) to a GPO generation script (<https://github.com/jzadeh/Aktaion/tree/master/python>)

# Aktaion Logical Workflow



# Learning = Compression?

- Train a model to predict mpg as a function of car weight, number of cylinders and displacement

	mtcars										
	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1

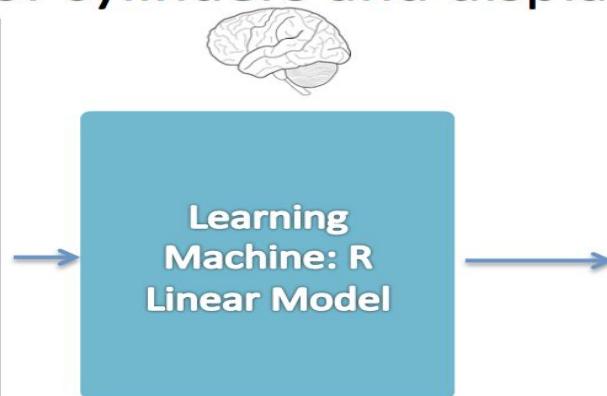


Learning  
Machine: R  
Linear Model

# Learning = Compression?

- Train a model to predict mpg as a function of car weight, number of cylinders and displacement

```
> mtcars
   mpg cyl disp hp drat wt qsec vs am gear carb
Mazda RX4       21.0   6 160.0 110 3.90 2.620 16.46  0  1   4   4
Mazda RX4 Wag   21.0   6 160.0 110 3.90 2.875 17.02  0  1   4   4
Datsun 710      22.8   4 108.0  93 3.85 2.320 18.61  1  1   4   1
Hornet 4 Drive   21.4   6 258.0 110 3.08 3.215 19.44  1  0   3   1
Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02  0  0   3   2
Valiant        18.1   6 225.0 105 2.76 3.460 20.22  1  0   3   1
Duster 360      14.3   8 360.0 245 3.21 3.570 15.84  0  0   3   4
Merc 240D       24.4   4 146.7  62 3.69 3.190 20.00  1  0   4   2
Merc 230        22.8   4 140.8  95 3.92 3.150 22.90  1  0   4   2
Merc 280        19.2   6 167.6 123 3.92 3.440 18.30  1  0   4   4
Merc 280C       17.8   6 167.6 123 3.92 3.440 18.90  1  0   4   4
Merc 450SE      16.4   8 275.8 180 3.07 4.070 17.40  0  0   3   3
Merc 450SL      17.3   8 275.8 180 3.07 3.730 17.60  0  0   3   3
Merc 450SLC     15.2   8 275.8 180 3.07 3.780 18.00  0  0   3   3
Cadillac Fleetwood 10.4   8 472.0 205 2.93 5.250 17.98  0  0   3   4
Lincoln Continental 10.4   8 460.0 215 3.00 5.424 17.82  0  0   3   4
Chrysler Imperial 14.7   8 440.0 230 3.23 5.345 17.42  0  0   3   4
Fiat 128         32.4   4  78.7  66 4.08 2.200 19.47  1  1   4   1
Honda Civic      30.4   4  75.7  52 4.93 1.615 18.52  1  1   4   2
Toyota Corolla    33.9   4  71.1  65 4.22 1.835 19.90  1  1   4   1
Toyota Corona    21.5   4 120.1  97 3.70 2.465 20.01  1  0   3   1
```



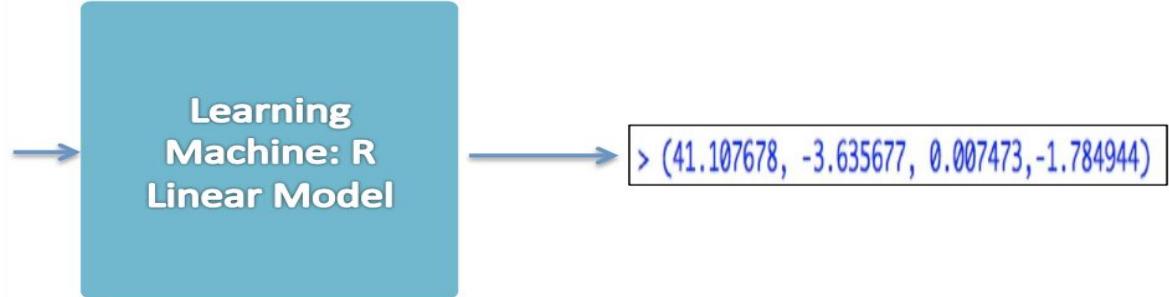
```
> mfit = lm(mpg ~ wt + disp + cyl, data=mtcars)
```

# Learning = Compression?

- The overall input data is reduced in a “compressed form” to use in future predictions

```
> mtcars
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1



```
> mfit = lm(mpg ~ wt + disp + cyl, data=mtcars)
```

# **SECURITY ANALYTICS: THE LAMBDA DEFENSE**

# Operator Time is Valuable!

- Googles Experience with ML in Cybersecurity:  
<https://web.stanford.edu/class/cs259d/lectures/Session11.pdf>

Google

## Why is statistical anomaly detection for security hard?

	Learning	Cost of error FP / FN	Goal	Attacker
Anomaly detection	Unsupervised	Medium / High	Classify & explain	Adaptive
Spam detection	Supervised	High / Low	Classify	Adaptive
Product recommendation	Supervised	Low / Low	Classify	N/A

For more on this topic, consider “On Using Machine Learning For Network Intrusion Detection”, Sommer and Paxson, Oakland 2010.

Google

## An alternative: rule-based detection

Manually created rules  
→ Characterize attacks or deviations, not normalcy

- Locality: each rule covers a small set of logs and features
- Explainability: direct connection between rules and alerts
- Specificity: make better use of analysts' expertise

# Lambda Architecture

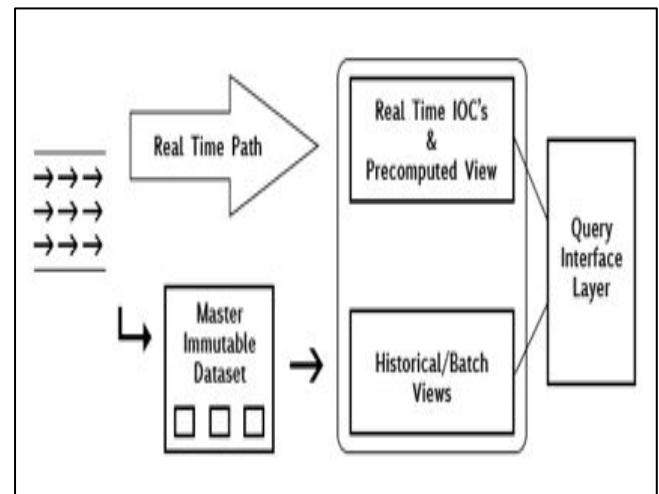
- The reference to the pattern comes from “Lambda Architecture” Coined by Nathan Marz
- GOTO 2012 • Runaway Complexity in Big Data Systems...and a Plan to Stop it • Nathan Marz: <https://www.youtube.com/watch?v=ucHjyb6jvo8>

Architecture is described by three simple equations:

**batch view = function(all data)**

**realtime view = function(realtime view, new data)**

**query = function(batch view, realtime view)**



# The Lambda Defense: Decomposing Behaviors for Intrusion Detection

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .
4. Behavioral Model: For each sub component  $B_i, NB_i$  we assign one or more models  $M_1(B_i), \dots, M_k(B_i), M_{k+1}(NB_i), M_{k+2}(NB_i), \dots$  and we evolve these models over time in the next step. We keep the definition of model intentionally vague here but in practice each model is usually machine learning based (supervised/semi-supervised/unsupervised) or heuristically engineered around a use case that is not driven by a pattern recognition problem. It is important to note in that we are looking for the right tool for the job and machine learning is sometimes only used as a pre-filtering step or not at all. In some cases we combine graph processing or techniques from signal analysis where it helps with overall problem of a specific model implementation.
5. Map/Reduce Based Model Life Cycle: We can update simultaneous models in a scalable way with a simple map and reduce operation given by a signature of the form [41]:  $ModelRDD.aggregate[M](seqOp: (M, Data) \Rightarrow M(Data), combOp: (M_{t_0}, M_{t_1}) \Rightarrow M_{t_2})$

# Step 1: Break the problem into use cases

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.

# Step 2: Find what use cases have highest security impact

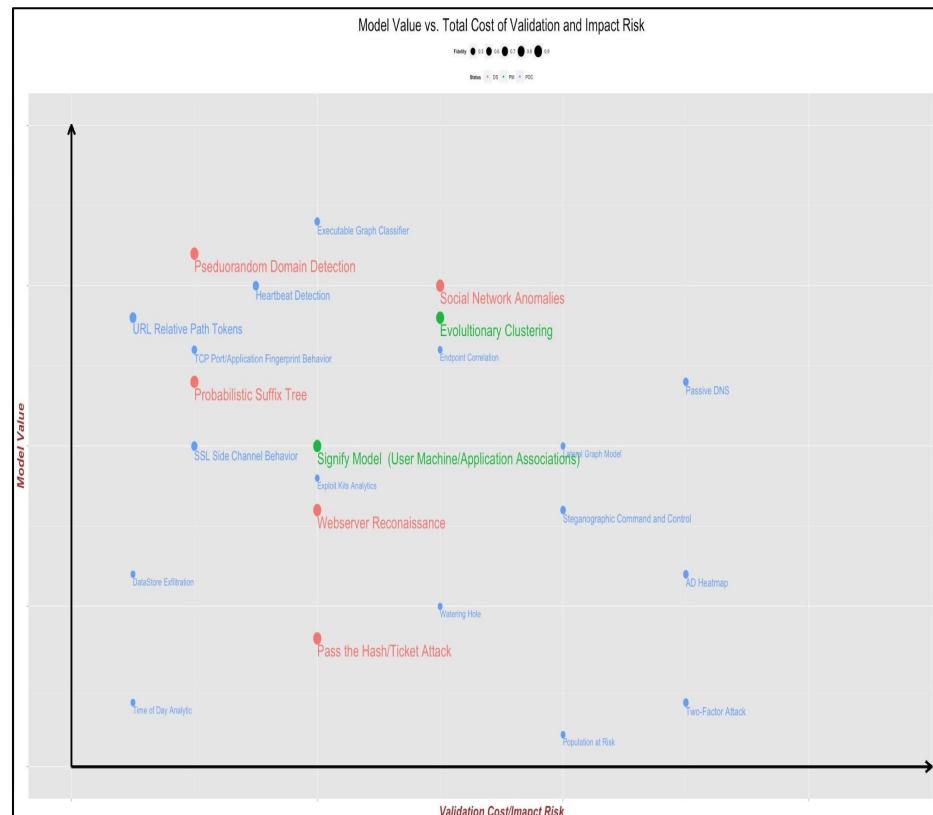
## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).

# Step 2: Find what use cases have highest security impact

## DEFINITION 1. *The Lambda Defense Pattern*

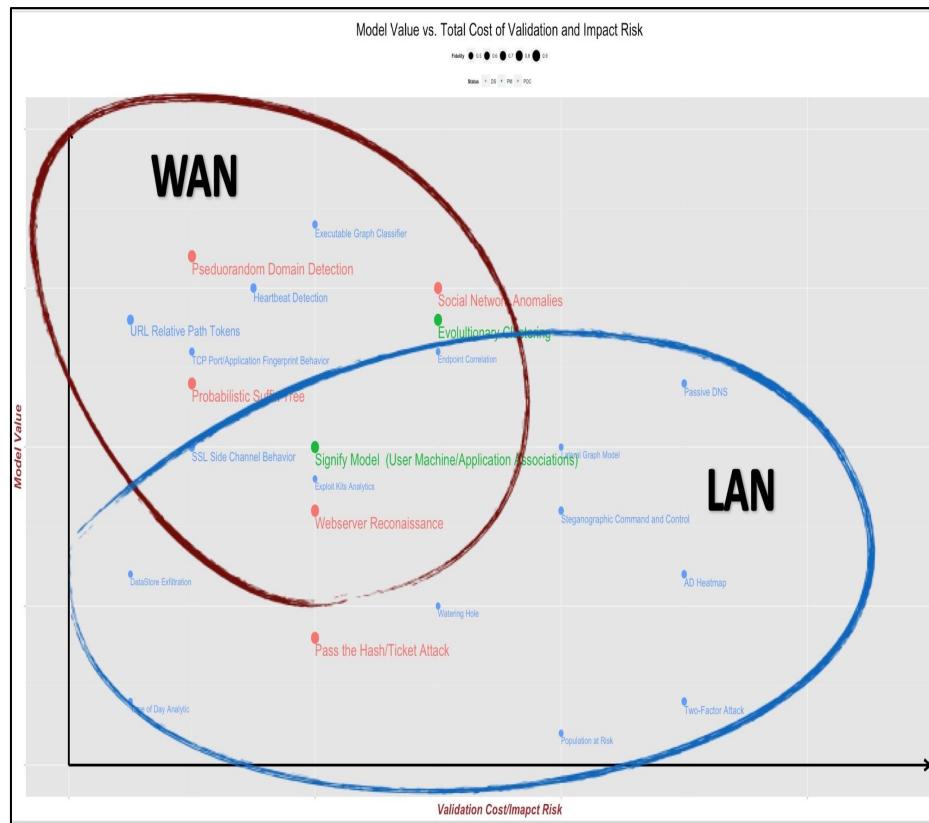
1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).



# Step 2: Find what use cases have highest security impact

## DEFINITION 1. *The Lambda Defense Pattern*

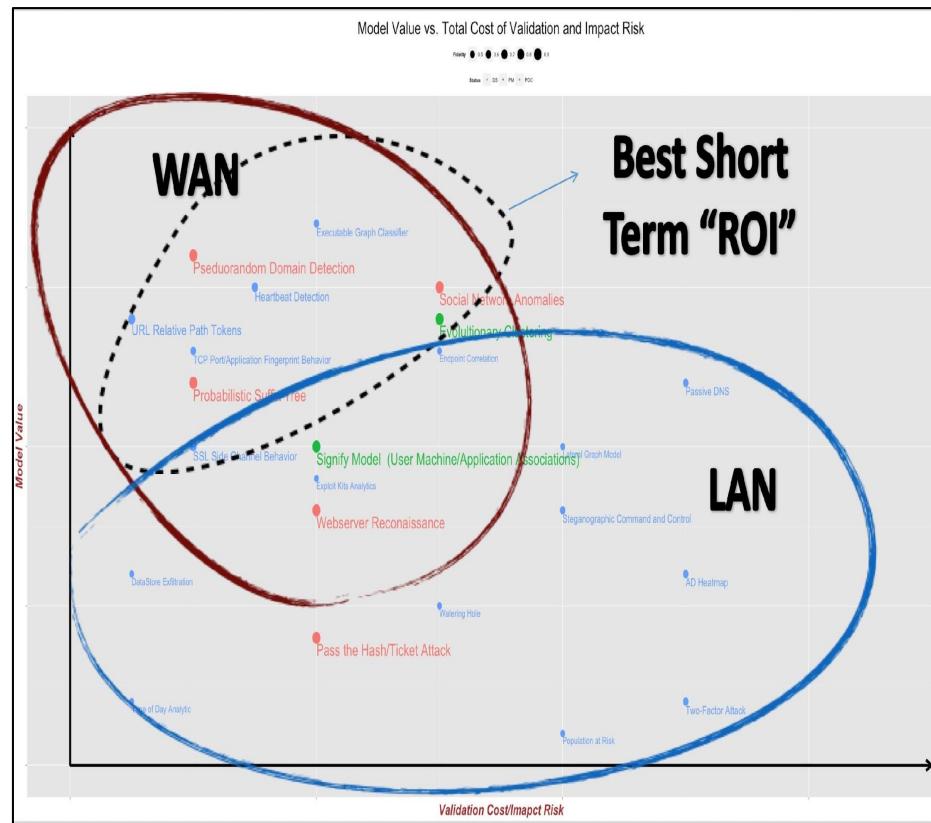
1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).



# Step 2: Find what use cases have highest security impact

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).



# Step 3: Decompose the problem into two types of computation

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

# Step 3: Decompose the problem into two types of computation

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

**Arbitrary User Behavior = Sequential Component + “Un-Ordered” Component**

# Step 3: Decompose the problem into two types of computation

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

**Arbitrary User Behavior = Sequential Component + “Un-Ordered” Component**

## Examples

- Sequential Behaviors
  1. Exploit Chains
  2. Timing Analysis (Periodicity)
  3. Active Directory Sequence
  4. Authentication Graph
- Non Sequential Behaviors
  1. Fingerprinting
  2. Grouping Behaviors
  3. Application Counts
  4. Rare file extension counts for Webshell detection

# Step 3: Decompose the problem into two types of computation

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

**Arbitrary User Behavior = Sequential Component + “Un-Ordered” Component**

## Mapping Behaviors to Computational Paths

- Easy to Parallelize
  1. Count()
  2. Average()
  3. Time series()
  4. Local state computations
    - Per user/IP/account/...
- Hard to Parallelize (NC Complete Complexity)
  1. Rank()
  2. Median
  3. Anything that keeps track of global state
  4. Machine Learning Computations
    - Matrix Inversion or
    - Stochastic Gradient Descent

# Step 3: Optimizing the use of technology for problem complexity

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

**Arbitrary User Behavior = Sequential Component + “Un-Ordered” Component**

## Mapping Behaviors to Computational Paths

- Easy to Parallelize
  1. Count()
  2. Average()
  3. Time series()
  4. Local state computations
    - Per user/IP/account/...
- Hard to Parallelize (NC Complete Complexity)
  1. Rank()
  2. Median
  3. Anything that keeps track of global state
  4. Machine Learning Computations
    - Matrix Inversion or
    - Stochastic Gradient Descent

In complexity theory, the notion of **P-complete** decision problems is useful in the analysis of both:

1. which problems are difficult to parallelize effectively, and;
2. which problems are difficult to solve in limited space.

# Step 4: Build an ML Model for each important sub-behavior

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

4. Behavioral Model: For each sub component  $B_i, NB_i$  we assign one or more models  $M_1(B_i), \dots, M_k(B_i), M_{k+1}(NB_i), M_{k+2}(NB_i), \dots$  and we evolve these models over time in the next step. We keep the definition of model intentionally vague here but in practice each model is usually machine learning based (supervised/semi-supervised/unsupervised) or heuristically engineered around a use case that is not driven by a pattern recognition problem. It is important to note in that we are looking for the right tool for the job and machine learning is sometimes only used as a pre-filtering step or not at all. In some cases we combine graph processing or techniques from signal analysis where it helps with overall problem of a specific model implementation.

# Step 4: Build an ML Model for each important sub-behavior

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally

speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .

4. Behavioral Model: For each sub component  $B_i, NB_i$  we assign one or more models  $M_1(B_i), \dots, M_k(B_i), M_{k+1}(NB_i), M_{k+2}(NB_i), \dots$  and we evolve these models over time in the next step. We keep the definition of model intentionally vague here but in practice each model is usually machine learning based (supervised/semi-supervised/unsupervised) or heuristically engineered around a use case that is not driven by a pattern recognition problem. It is important to note in that we are looking for the right tool for the job and machine learning is sometimes only used as a pre-filtering step or not at all. In some cases we combine graph processing or techniques from signal analysis where it helps with overall problem of a specific model implementation.

Each Model can be batch, real-time or hybrid mode

Domain Name	TotalCnt	RiskFactor	AGD	SessionTime	RefEntropy	Nulls
evcertsolutions.org	144	6.05	1	1	0	0
jeveyd37w9.com	6192	5.05	0	1	0	0
cdtbs.stevehousemedia.com	107	3	0	0	0	0
log.tatrade.com	111	2	0	1	0	0
go.vdrprocess.com	170	2	0	0	0	0
state.weternridgeline.com	310	2	0	1	0	0
cdtbs.stevehousemedia.com	107	1	0	0	0	0
log.tatrade.com	111	1	0	1	0	0

```
//decorate the (domain,feature vector) pair with predicted output of the model
//use pair RDD's to recover the full data we need for anomaly generation by a inner join
val predictedOutput = vectorizeFeatures.map(x => (x._1, randomForestModel.predict(x._2)))
```

# Step 5: Operationalize the Model Life Cycle

## DEFINITION 1. *The Lambda Defense Pattern*

1.  $|S| = |\mathbb{R}|$ : Start with the threat surface  $S$  for a fixed adversarial environment. (The examples in this paper  $S$  will usually be the threat surface associated to a large enterprise network). The space of possible behaviors is a continuum at this point.
2. Enumerate the threat surface: How do we reduce the space of possible behaviors to actionable units that we can process in an upstream functional workflow? To answer this question we project down to modeling only a finite set  $S \approx \{TTP_1, \dots, TTP_n\}$ . For a given threat surface the art of the initial design is to identify the core attack patterns needed to be modeled and to anticipate which of those parts are most likely to “drift” [19]. Better yet assign metrics to this finite approximation of the threat surface (see figure 1).
3. Functional Decomposition: For each  $TTP_i$  we model the behavior as  $TTP_i := B_i + NB_i$ . Where  $B_i$  is what we call the sequential component of the behavior and  $NB_i$  is called the non-sequential component. Formally speaking each  $TTP_i$  is actually a stochastic process (sequence of random variable)  $TTP_i = \{TTP_i(t) : t \geq 0\}$ .
4. Behavioral Model: For each sub component  $B_i, NB_i$  we assign one or more models  $M_1(B_i), \dots, M_k(B_i), M_{k+1}(NB_i), M_{k+2}(NB_i), \dots$  and we evolve these models over time in the next step. We keep the definition of model intentionally vague here but in practice each model is usually machine learning based (supervised/semi-supervised/unsupervised) or heuristically engineered around a use case that is not driven by a pattern recognition problem. It is important to note in that we are looking for the right tool for the job and machine learning is sometimes only used as a pre-filtering step or not at all. In some cases we combine graph processing or techniques from signal analysis where it helps with overall problem of a specific model implementation.
5. Map/Reduce Based Model Life Cycle: We can update simultaneous models in a scalable way with a simple map and reduce operation given by a signature of the form [41]:  $ModelRDD.aggregate[M](seqOp: (M, Data) \Rightarrow M(Data), combOp: (M_{t_0}, M_{t_1}) \Rightarrow M_{t_2})$

**Q+A**

# The Tool

IDS Source Code: <http://www.github.com/jzadeh/chiron-elk>

VM ISO: [https://drive.google.com/uc?id=0B8G0yrax\\_wtaNzAxM2ptdUZ6anc&export=download](https://drive.google.com/uc?id=0B8G0yrax_wtaNzAxM2ptdUZ6anc&export=download)

## Requirements

- Can be deployed on a VM - 60GB HD 1GB RAM
- Chiron is good for a home network 5 - 10 users
- 30 day log retention standard